

To study the smoking effects on metabolic health and Inflammation Markers

Objectives:

To build a model which predicts diabetic category using smokers and non smokers body signal

Comparing the features of smokers and non smokers using exploratory data analysis.

To study the association of smoking on blood sugar level and dental caries.

Determining whether smokers cholesterol level influencing on blood pressure.

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

ID : serial number given to a person.

gender: the gender of a person being either female(0) or male(1).

age : 5-years gap age groups.

height(cm)

weight(kg)

waist(cm) : Waist circumference length

eyesight(left): Denoted by a value between 0.1-2.5, or 9.9

eyesight(right): Denoted by a value between 0.1-2.5, or 9.9

hearing(left): hearing of the person's ear

hearing(right): hearing of the person's ear

systolic : blood pressure.

relaxation : blood pressure.

fasting blood sugar: blood sugar before meals level

Cholesterol : total

triglyceride

HDL : cholesterol type

LDL : cholesterol type

hemoglobin

serum creatinine

AST : glutamic oxaloacetic transaminase type

ALT : glutamic oxaloacetic transaminase type

Gtp : γ -GTP

oral : Oral Examination status (s whether the examinee accepted the oral examination).

dental caries

tartar : tartar status

smoking: smoking status of a person

In [2]:

```
1 df=pd.read_csv("smoking.csv")
2 df
```

Out[2]:

	ID	smoking	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
0	0	No	F	40	155	60	81.3	1.2	
1	1	No	F	40	160	60	81.0	0.8	
2	2	Yes	M	55	170	60	80.0	0.8	
3	3	No	M	40	165	70	88.0	1.5	
4	4	No	F	40	155	60	86.0	1.0	
...
55687	55676	No	F	40	170	65	75.0	0.9	
55688	55681	No	F	45	160	50	70.0	1.2	
55689	55683	No	F	55	160	50	68.5	1.0	
55690	55684	No	M	60	165	60	78.0	0.8	
55691	55691	Yes	M	55	160	65	85.0	0.9	

55692 rows × 26 columns

In [3]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55692 entries, 0 to 55691
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                    55692 non-null  int64
1   smoking                             55692 non-null  object
2   gender                              55692 non-null  object
3   age                                  55692 non-null  int64
4   height(cm)                          55692 non-null  int64
5   weight(kg)                          55692 non-null  int64
6   waist(cm)                           55692 non-null  float64
7   eyesight(left)                      55692 non-null  float64
8   eyesight(right)                    55692 non-null  float64
9   hearing(left)                       55692 non-null  object
10  hearing(right)                      55692 non-null  object
11  systolic                            55692 non-null  int64
12  relaxation                          55692 non-null  int64
13  fasting blood sugar                 55692 non-null  int64
14  Cholesterol                         55692 non-null  int64
15  triglyceride                        55692 non-null  int64
16  HDL                                 55692 non-null  int64
17  LDL                                 55692 non-null  int64
18  hemoglobin                          55692 non-null  float64
19  serum creatinine                    55692 non-null  float64
20  AST                                 55692 non-null  int64
21  ALT                                 55692 non-null  int64
22  Gtp                                 55692 non-null  int64
23  oral                                55692 non-null  object
24  dental caries                       55692 non-null  int64
25  tartar                              55692 non-null  object
dtypes: float64(5), int64(15), object(6)
memory usage: 11.0+ MB
```

```
In [4]: 1 df.isnull().sum()
```

```
Out[4]: ID                                0
smoking                                0
gender                                0
age                                    0
height(cm)                            0
weight(kg)                            0
waist(cm)                             0
eyesight(left)                        0
eyesight(right)                       0
hearing(left)                         0
hearing(right)                        0
systolic                              0
relaxation                            0
fasting blood sugar                   0
Cholesterol                           0
triglyceride                          0
HDL                                    0
LDL                                    0
hemoglobin                            0
serum creatinine                      0
AST                                    0
ALT                                    0
Gtp                                    0
oral                                  0
dental caries                         0
tartar                               0
dtype: int64
```

```
In [5]: 1 df.columns
```

```
Out[5]: Index(['ID', 'smoking', 'gender', 'age', 'height(cm)', 'weight(kg)',
              'waist(cm)', 'eyesight(left)', 'eyesight(right)', 'hearing(left)',
              'hearing(right)', 'systolic', 'relaxation', 'fasting blood sugar',
              'Cholesterol', 'triglyceride', 'HDL', 'LDL', 'hemoglobin',
              'serum creatinine', 'AST', 'ALT', 'Gtp', 'oral', 'dental caries',
              'tartar'],
              dtype='object')
```

```
In [6]: 1 def sys_cat(sys_bp):
2         if sys_bp>=0 and sys_bp<=90:
3             return "Hypotension"
4         elif sys_bp>=90 and sys_bp<=120:
5             return "Normal"
6         elif sys_bp>=120 and sys_bp<=130:
7             return "Elevated"
8         elif sys_bp>=130 and sys_bp<=140:
9             return "Hypertension stage1"
10        elif sys_bp>=140 and sys_bp<=150:
11            return "Hypertension stage2"
12        else:
13            return "Hypertension crisis"
14
15 df['systolic_category']=df['systolic'].apply(sys_cat)
```

```
In [7]: 1 def rel_cat(rel_bp):
2         if rel_bp>=0 and rel_bp<=60:
3             return "Hypotension"
4         elif rel_bp>=60 and rel_bp<=80:
5             return "Normal"
6         elif rel_bp>=80 and rel_bp<=90:
7             return "Elevated"
8         elif rel_bp>=90 and rel_bp<=100:
9             return "Hypertension stage1"
10        elif rel_bp>=100 and rel_bp<=120:
11            return "Hypertension stage2"
12        else:
13            return "Hypertension crisis"
14
15 df['relaxation_category']=df['relaxation'].apply(rel_cat)
```

```
In [8]: 1 def diab_cat(sugar):
2         if sugar>=0 and sugar<=80:
3             return "Hypoglycemia"
4         elif sugar>=80 and sugar<=100:
5             return "Normal"
6         elif sugar>=100 and sugar<=130:
7             return "Pre-Diabetic"
8         else:
9             return "Diabetic"
10
11 df['diabetic_category']=df['fasting blood sugar'].apply(diab_cat)
```

```
In [9]: 1 df['dental caries'] = df['dental caries'].replace({0: 'Yes', 1: 'No'})
2
```

In [10]:

```
1 df=df.loc[:,['ID', 'smoking', 'gender', 'age', 'height(cm)', 'weight(kg)',
2             'waist(cm)', 'eyesight(left)', 'eyesight(right)', 'hearing(left)',
3             'hearing(right)', 'systolic','systolic_category', 'relaxation','relaxation_category',
4             'diabetic_category','Cholesterol', 'triglyceride', 'HDL', 'LDL',
5             'serum creatinine', 'AST', 'ALT', 'Gtp', 'oral', 'dental caries', 'tartar']]
6
7 df
```

Out[10]:

	ID	smoking	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
	0	0	No	F	40	155	60	81.3	1.2
	1	1	No	F	40	160	60	81.0	0.8
	2	2	Yes	M	55	170	60	80.0	0.8
	3	3	No	M	40	165	70	88.0	1.5
	4	4	No	F	40	155	60	86.0	1.0

	55687	55676	No	F	40	170	65	75.0	0.9
	55688	55681	No	F	45	160	50	70.0	1.2
	55689	55683	No	F	55	160	50	68.5	1.0
	55690	55684	No	M	60	165	60	78.0	0.8
	55691	55691	Yes	M	55	160	65	85.0	0.9

55692 rows × 29 columns

In []:

1

In []:

1

Exploratory Data Analysis

In [11]:

```
1 df["smoking"].value_counts()
```

Out[11]:

No	35237
Yes	20455

Name: smoking, dtype: int64

In [12]:

```
1 df["smoking"].value_counts()
```

Out[12]:


No	35237
Yes	20455

Name: smoking, dtype: int64

Total number of smokers and non-smokers

In [13]:

```
1 df["smoking"].value_counts().plot.pie(shadow=True,explode=[0,0.1],autopct='%1.1f%%')
2 plt.title("Number of Smokers and Non-smokers")
3 plt.show()
```




In [14]:

```
1 # only 36.7% people smoke
2 # 63.3% people do not smoke
```

In [15]:

```
1 df["gender"].value_counts().plot.pie(shadow=True,explode=[0,0.1],autopct='%1.1f%%')
2 plt.title("Number of Male and Female")
3 plt.show()
```



In [16]:

```
1 # Total 63.6% are male
2 # 36.4% are female
```

In [17]:

```
1 smokeYes=df[df["smoking"]=="Yes"]
2 smokeYes.head(2)
```

Out[17]:

	ID	smoking	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
2	2	Yes	M	55	170	60	80.0	0.8	0.8
6	6	Yes	M	40	160	60	85.5	1.0	1.0


2 rows × 29 columns

```
In [18]: 1 smokeYes['gender'].value_counts()
```

```
Out[18]: M    19596  
         F     859  
         Name: gender, dtype: int64
```

```
In [19]: 1 smokeYes["gender"].value_counts().plot.pie(shadow=True,explode=[0,0.1],  
2 plt.title("Number of male and female smokers")  
3 plt.show()  
4
```

Number of male and female smokers



Average age of male and female smokers.

```
In [20]: 1 smokeNo=df[df["smoking"]=="No"]  
2 #smokeNo
```

```
In [21]: 1 smokeYesW=smokeYes[smokeYes["gender"]=="F"]  
2 #smokeYesW
```


In [22]:

1

smokeYesM=smokeYes[smokeYes["gender"]=="M"]

2

smokeYesM

Out[22]:

	ID	smoking	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
	2	2	Yes	M	55	170	60	80.0	0.8
	6	6	Yes	M	40	160	60	85.5	1.0
	12	13	Yes	M	35	170	70	81.0	1.5
	17	19	Yes	M	35	165	70	87.5	1.0
	18	21	Yes	M	60	165	65	79.0	1.0

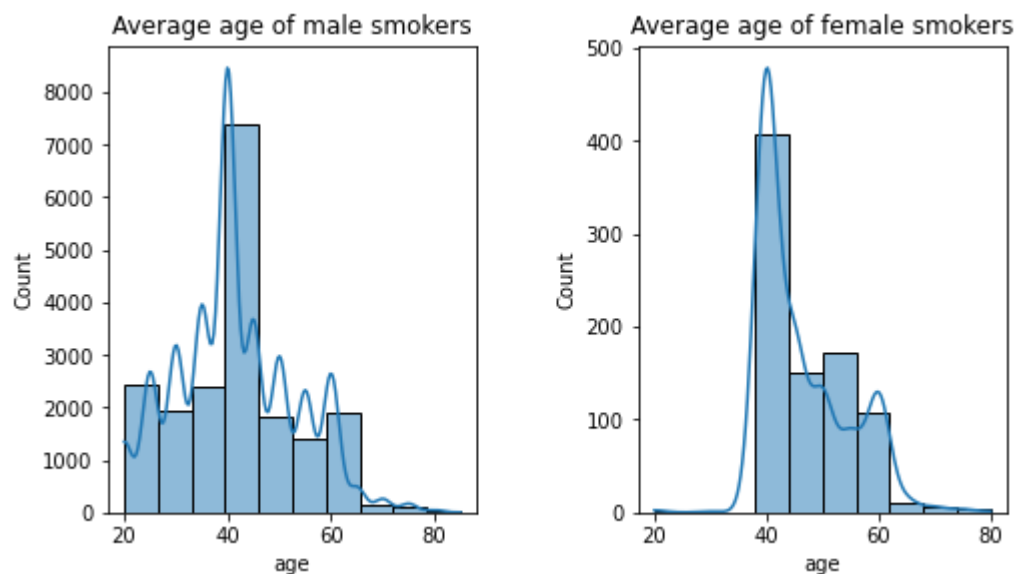
	55673	55596	Yes	M	75	155	45	68.0	0.6
	55675	55607	Yes	M	30	170	70	88.8	1.0
	55679	55632	Yes	M	45	180	90	92.0	1.2
	55684	55666	Yes	M	40	170	65	85.0	1.2
	55691	55691	Yes	M	55	160	65	85.0	0.9

19596 rows × 29 columns

```
In [23]: 1 plt.figure(figsize=(8,5))
2
3 plt.subplot(1,2,1)
4 sns.histplot(data=smokeYesM,x='age',kde=True,bins=10)
5 plt.title("Average age of male smokers")
6
7 plt.subplot(1,2,2)
8
9 sns.histplot(data=smokeYesW,x='age',kde=True,bins=10)
10 plt.title("Average age of female smokers")
11
12
13
14 plt.tight_layout(4)
15 plt.show()
```

C:\Users\darsh\AppData\Local\Temp\ipykernel_18136\3942839391.py:14: MatplotlibDeprecationWarning: Passing the pad parameter of tight_layout() positionally is deprecated since Matplotlib 3.3; the parameter will become keyword-only two minor releases later.

```
plt.tight_layout(4)
```



```
In [24]: 1 w=sum(smokeYesW['age'])/len(smokeYesW['age'])
2 print('Average age of female smokers',w)
```

Average age of female smokers 46.38533178114086

```
In [25]: 1 m=sum(smokeYesM['age'])/len(smokeYesM['age'])
2 print('Average age male smokers',m)
```

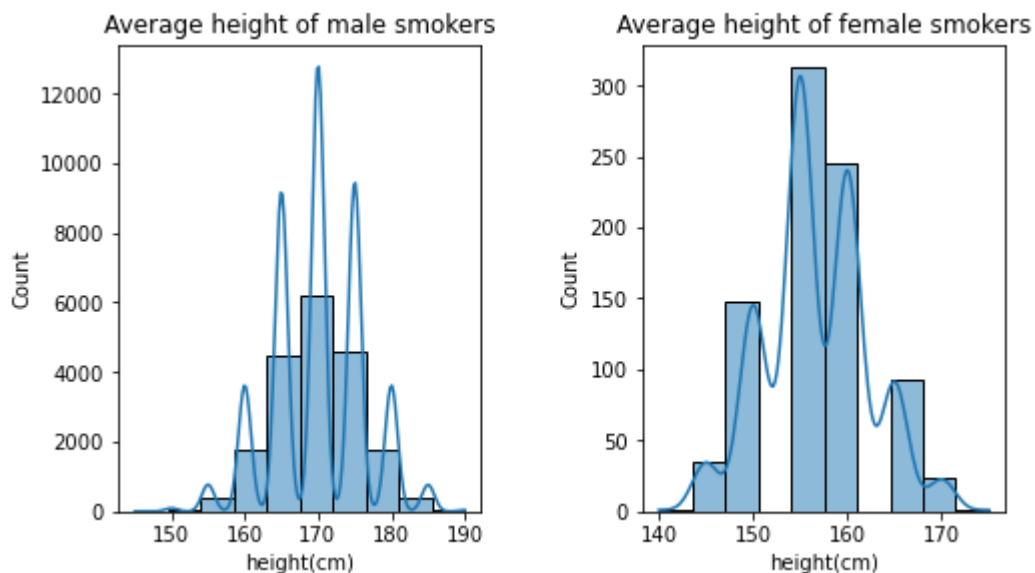
Average age male smokers 41.39798938558889

Average height of male and female smokers.

```
In [26]: 1 plt.figure(figsize=(8,5))
2
3 plt.subplot(1,2,1)
4 sns.histplot(data=smokeYesM,x='height(cm)',kde=True,bins=10)
5 plt.title("Average height of male smokers")
6
7
8
9 plt.subplot(1,2,2)
10 sns.histplot(data=smokeYesW,x='height(cm)',kde=True,bins=10)
11 plt.title("Average height of female smokers")
12
13 plt.tight_layout(4)
14 plt.show()
```

C:\Users\darsh\AppData\Local\Temp\ipykernel_18136\1465316697.py:13: MatplotlibDeprecationWarning: Passing the pad parameter of tight_layout() positionally is deprecated since Matplotlib 3.3; the parameter will become keyword-only two minor releases later.

```
plt.tight_layout(4)
```



```
In [27]: 1 w=sum(smokeYesW['height(cm)'])/len(smokeYesW['height(cm)'])
2 print('Average height of female smokers',w)
```

Average height of female smokers 156.64726426076834

```
In [28]: 1 m=sum(smokeYesM['height(cm)'])/len(smokeYesM['height(cm)'])
2 print('Average Height of male smokers',m)
```

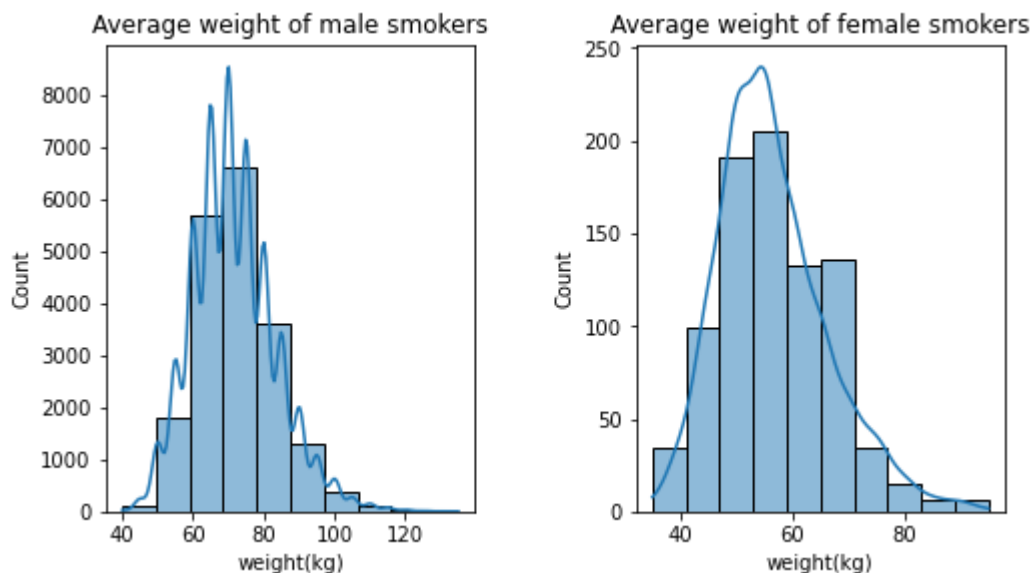
Average Height of male smokers 169.996938150643

Average weight male and female smokers

```
In [29]: 1 plt.figure(figsize=(8,5))
2
3 plt.subplot(1,2,1)
4 sns.histplot(data=smokeYesM,x='weight(kg)',kde=True,bins=10)
5 plt.title("Average weight of male smokers")
6
7 plt.subplot(1,2,2)
8 sns.histplot(data=smokeYesW,x='weight(kg)',kde=True,bins=10)
9 plt.title("Average weight of female smokers")
10
11
12
13 plt.tight_layout(4)
14 plt.show()
```

C:\Users\darsh\AppData\Local\Temp\ipykernel_18136\2753398324.py:13: MatplotlibDeprecationWarning: Passing the pad parameter of tight_layout() positionally is deprecated since Matplotlib 3.3; the parameter will become keyword-only two minor releases later.

```
plt.tight_layout(4)
```




```
In [30]: 1 w=sum(smokeYesW['weight(kg)'])/len(smokeYesW['weight(kg)'])
2 print('Average weight of female smokers',w)
```

Average weight of female smokers 56.44935972060536

```
In [31]: 1 m=sum(smokeYesM['weight(kg)'])/len(smokeYesM['weight(kg)'])
2 print('Average weight of female smokers',m)
```


Average weight of female smokers 71.59573382322922

```
In [32]: 1 plt.figure(figsize=(10,10))
2
3 plt.boxplot([smokeYesM['age'],smokeYesM['height(cm)'], smokeYesM['weight']
4 plt.title('Box Plot of Age, Height, and Weight of Male Smokers')
5 plt.ylabel('Count')
6 plt.grid(True)
7 plt.show()
```



In [33]:

```
1 plt.figure(figsize=(10,10))
2
3 plt.boxplot([smokeYesW['age'],smokeYesW['height(cm)'], smokeYesW['weight(kg)'])
4 plt.title('Box Plot of Age, Height, and Weight of Female Smokers')
5 plt.ylabel('Count')
6 plt.grid(True)
7 plt.show()
```



In [34]:

```
1 smokeYes.head()
```

Out[34]:


	ID	smoking	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
2	2	Yes	M	55	170	60	80.0	0.8	0.8
6	6	Yes	M	40	160	60	85.5	1.0	1.0
12	13	Yes	M	35	170	70	81.0	1.5	1.0
17	19	Yes	M	35	165	70	87.5	1.0	0.8
18	21	Yes	M	60	165	65	79.0	1.0	1.0

5 rows × 29 columns

```
In [35]: 1 plt.figure(figsize=(8,8))
2 plt.subplot(1,2,1)
3
4 smokeYes["hearing(left)"].value_counts().plot.pie(shadow=True,explode=[0.05,0.05])
5
6
7 plt.subplot(1,2,2)
8 smokeYes["hearing(right)"].value_counts().plot.pie(shadow=True,explode=[0.05,0.05])
9
10 plt.tight_layout(4)
11 plt.show()
```

C:\Users\darsh\AppData\Local\Temp\ipykernel_18136\2654710606.py:10: MatplotlibDeprecationWarning: Passing the pad parameter of tight_layout() positionally is deprecated since Matplotlib 3.3; the parameter will become keyword-only two minor releases later.

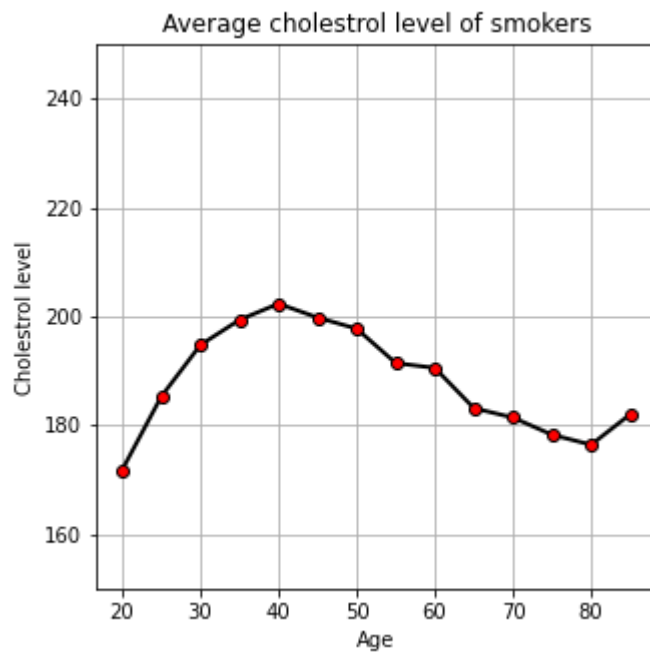
plt.tight_layout(4)



Average cholestrol level of smokers


In [36]:

```
1 smokeYes.groupby('age').mean()['Cholesterol'].plot(ylim=(150,250),color='red',  
2 markerfacecolor='red',  
3  
4 plt.xlabel('Age')  
5 plt.ylabel("Cholestrol level")  
6 plt.title("Average cholestrol level of smokers")  
7 plt.grid()  
8 plt.show()  
9  
10
```



Average cholestrol level of non-smokers


```
In [37]: 1 smokeNo.groupby('age').mean()['Cholesterol'].plot(ylim=(150,250),color=
2 markerfacecolor='red')
3 plt.xlabel('Age')
4 plt.ylabel("Cholestrol level")
5 plt.title("Average cholestrol level of Non-smokers")
6 plt.grid()
7
8
9 plt.show()
```



```
In [38]: 1 smokeyes=df[df["smoking"]=="Yes"]
2 smokeyes.head(2)
```

Out[38]:

	ID	smoking	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
2	2	Yes	M	55	170	60	80.0	0.8	0.8
6	6	Yes	M	40	160	60	85.5	1.0	1.0

2 rows × 29 columns


```
In [39]: 1 smokeno=df[df["smoking"]=="No"]
2 smokeno.head(2)
```

Out[39]:

	ID	smoking	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
0	0	No	F	40	155	60	81.3	1.2	1.0
1	1	No	F	40	160	60	81.0	0.8	0.6

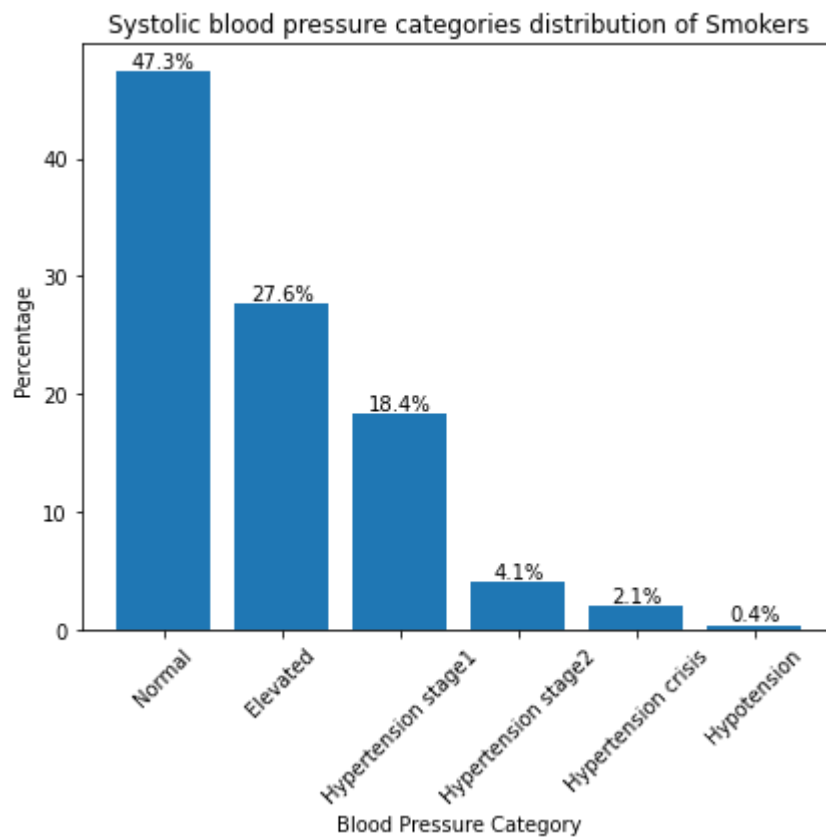
2 rows × 29 columns

```
In [40]: 1 plt.figure(figsize=(12,12))
2
3 counts = smokeyes['systolic_category'].value_counts()
4 counts1= smokeno['systolic_category'].value_counts()
5
6 percentages = (counts / counts.sum()) * 100
7 percentages1 = (counts1 / counts1.sum()) * 100
8
9 sorted_categories = percentages.sort_values(ascending=False).index
10 sorted_categories1 = percentages1.sort_values(ascending=False).index
11
12
13 plt.subplot(1,2,1)
14
15 plt.bar(sorted_categories, percentages[sorted_categories])
16
17 plt.xlabel('Blood Pressure Category')
18 plt.ylabel('Percentage')
19 plt.title('Systolic blood pressure categories distribution of Smokers')
20 plt.xticks(rotation=45)
21
22 for i, v in enumerate(percentages[sorted_categories]):
23     plt.text(i, v, f'{v:.1f}%', ha='center', va='bottom')
24
25
26
27 plt.subplot(1,2,2)
28
29 plt.bar(sorted_categories1, percentages1[sorted_categories1])
30
31 plt.xlabel('Blood Pressure Category')
32 plt.ylabel('Percentage')
33 plt.title('Systolic blood pressure categories distribution of Non smoker')
34 plt.xticks(rotation=45)
35
36 for i, v in enumerate(percentages1[sorted_categories1]):
37     plt.text(i, v, f'{v:.1f}%', ha='center', va='bottom')
38
39 plt.tight_layout()
40 plt.show()
41
```




Analysis of Systoli Blood Pressure Categories Distribution of Smokers

```
In [41]: 1 counts = smokeyes['systolic_category'].value_counts()
2
3
4 percentages = (counts / counts.sum()) * 100
5
6 sorted_categories = percentages.sort_values(ascending=False).index
7
8 plt.figure(figsize=(6,6))
9 plt.bar(sorted_categories, percentages[sorted_categories])
10
11 plt.xlabel('Blood Pressure Category')
12 plt.ylabel('Percentage')
13 plt.title('Systolic blood pressure categories distribution of Smokers')
14
15 plt.xticks(rotation=45)
16
17 for i, v in enumerate(percentages[sorted_categories]):
18     plt.text(i, v, f'{v:.1f}%', ha='center', va='bottom')
19
20 plt.tight_layout()
21 plt.show()
```




Analysis of Systolic Blood Pressure Categories Distribution of Non-smokers

```
In [42]: 1 counts = smokeno['systolic_category'].value_counts()
2
3
4 percentages = (counts / counts.sum()) * 100
5
6 sorted_categories = percentages.sort_values(ascending=False).index
7
8 plt.figure(figsize=(6,6))
9 plt.bar(sorted_categories, percentages[sorted_categories])
10
11 plt.xlabel('Blood Pressure Category')
12 plt.ylabel('Percentage')
13 plt.title('Blood Pressure Categories Distribution of Non-smokers')
14
15 plt.xticks(rotation=45)
16
17 for i, v in enumerate(percentages[sorted_categories]):
18     plt.text(i, v, f'{v:.1f}%', ha='center', va='bottom')
19
20 plt.tight_layout()
21 plt.show()
```




```
In [43]: 1 plt.figure(figsize=(12,10))
2
3 counts = smokeyes['relaxation_category'].value_counts()
4 counts1= smokeno['relaxation_category'].value_counts()
5
6 percentages = (counts / counts.sum()) * 100
7 percentages1 = (counts1 / counts1.sum()) * 100
8
9 sorted_categories = percentages.sort_values(ascending=False).index
10 sorted_categories1 = percentages1.sort_values(ascending=False).index
11
12
13 plt.subplot(1,2,1)
14
15 plt.bar(sorted_categories, percentages[sorted_categories])
16
17 plt.xlabel('Blood Pressure Category')
18 plt.ylabel('Percentage')
19 plt.title('Relaxation blood pressure categories distribution of Smokers')
20 plt.xticks(rotation=45)
21
22 for i, v in enumerate(percentages[sorted_categories]):
23     plt.text(i, v, f'{v:.1f}%', ha='center', va='bottom')
24
25
26
27 plt.subplot(1,2,2)
28
29 plt.bar(sorted_categories1, percentages1[sorted_categories1])
30
31 plt.xlabel('Blood Pressure Category')
32 plt.ylabel('Percentage')
33 plt.title('Relaxation blood pressure categories distribution of Non smokers')
34 plt.xticks(rotation=45)
35
36 for i, v in enumerate(percentages1[sorted_categories1]):
37     plt.text(i, v, f'{v:.1f}%', ha='center', va='bottom')
38
39 plt.tight_layout()
40 plt.show()
```



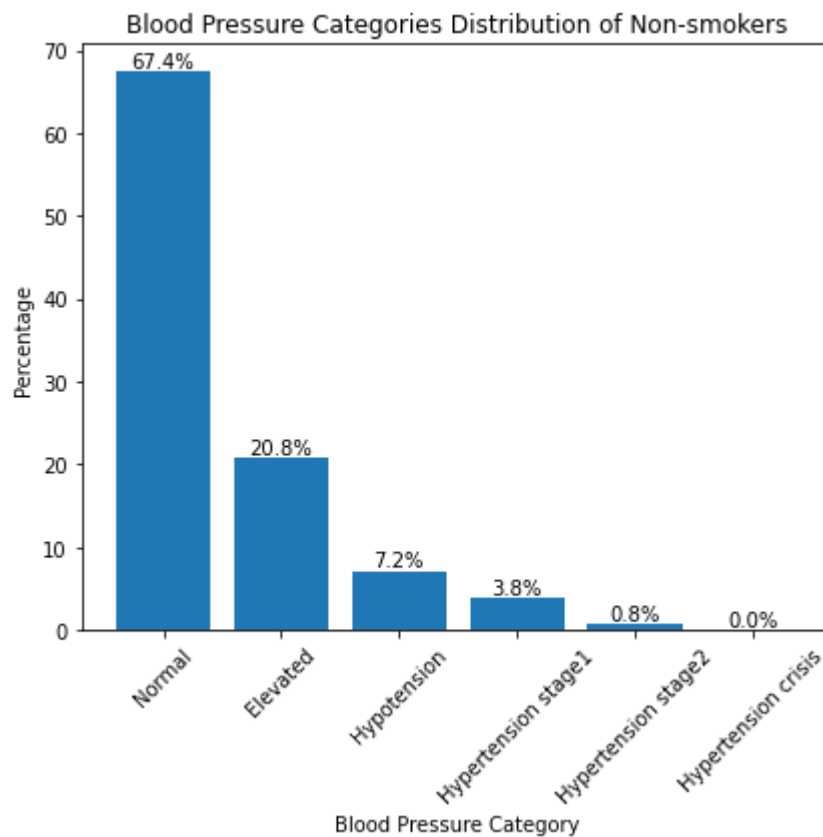
Analysis of Relaxation Blood Pressure Categories Distribution of smokers

```
In [44]: 1 counts = smokeyes['relaxation_category'].value_counts()
2
3
4 percentages = (counts / counts.sum()) * 100
5
6 sorted_categories = percentages.sort_values(ascending=False).index
7
8 plt.figure(figsize=(6,6))
9 plt.bar(sorted_categories, percentages[sorted_categories])
10
11 plt.xlabel('Blood Pressure Category')
12 plt.ylabel('Percentage')
13 plt.title('Blood Pressure Categories Distribution of Non-smokers')
14
15 plt.xticks(rotation=45)
16
17 for i, v in enumerate(percentages[sorted_categories]):
18     plt.text(i, v, f'{v:.1f}%', ha='center', va='bottom')
19
20 plt.tight_layout()
21 plt.show()
```



Analysis of Relaxation Blood Pressure Categories Distribution of Non-smokers


```
In [45]: 1 counts = smokeno['relaxation_category'].value_counts()
2
3
4 percentages = (counts / counts.sum()) * 100
5
6 sorted_categories = percentages.sort_values(ascending=False).index
7
8 plt.figure(figsize=(6,6))
9 plt.bar(sorted_categories, percentages[sorted_categories])
10
11 plt.xlabel('Blood Pressure Category')
12 plt.ylabel('Percentage')
13 plt.title('Blood Pressure Categories Distribution of Non-smokers')
14
15 plt.xticks(rotation=45)
16
17 for i, v in enumerate(percentages[sorted_categories]):
18     plt.text(i, v, f'{v:.1f}%', ha='center', va='bottom')
19
20 plt.tight_layout()
21 plt.show()
```




Pie chart of diabetes category of smokers and non-smokers

```
In [46]: 1 plt.figure(figsize=(10,10))
2 plt.subplot(1,2,1)
3
4 smokeYes["diabetic_category"].value_counts().plot.pie(shadow=True,autopct='%1.1f%%')
5 plt.title("Diabetes category of smokers")
6
7 plt.subplot(1,2,2)
8 smokeno["diabetic_category"].value_counts().plot.pie(shadow=True,autopct='%1.1f%%')
9
10 plt.title("Diabetes category of non-smokers")
11 plt.tight_layout(4)
12 plt.show()
```

C:\Users\darsh\AppData\Local\Temp\ipykernel_18136\3192570241.py:11: MatplotlibDeprecationWarning: Passing the pad parameter of tight_layout() positionally is deprecated since Matplotlib 3.3; the parameter will become keyword-only two minor releases later.

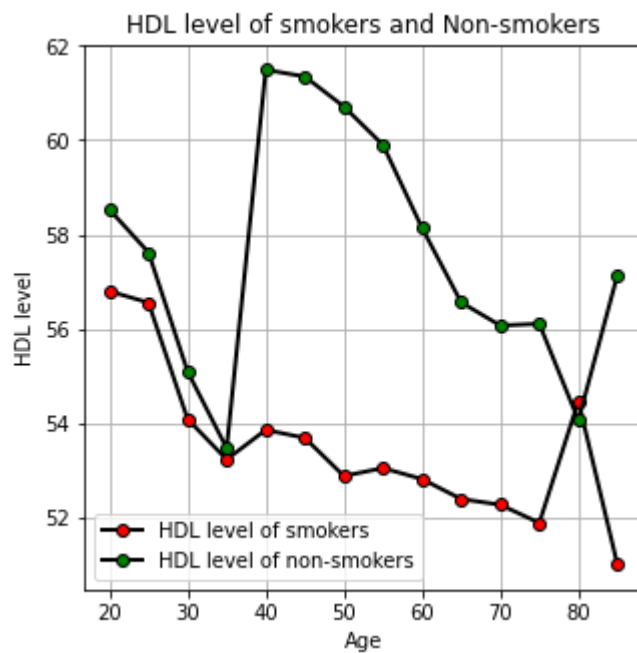
```
plt.tight_layout(4)
```



```
In [47]: 1 female=df[df['gender']=='F']
2 femaleYes=female[female['smoking']=='Yes']
3 femaleNo=female[female['smoking']=='No']
4
```


Analysis of average HDL level of smoker and non-smoker different age

```
In [48]: 1 plt.figure(figsize=(10,10))
2 smokeYes.groupby('age').mean()['HDL'].plot(color='k',marker='o',markerfacecolor='k',
3                                             linestyle='-',linewidth=2)
4 smokeNo.groupby('age').mean()['HDL'].plot(color='k',marker='o',markerfacecolor='k',
5                                             linestyle='-',linewidth=2)
6 plt.xlabel('Age')
7 plt.ylabel("HDL level")
8 plt.title("HDL level of smokers and Non-smokers")
9 plt.legend(loc='lower left')
10 plt.grid()
11 plt.show()
12
13
```



Analysis of average LDL level of smoker and non-smoker different age

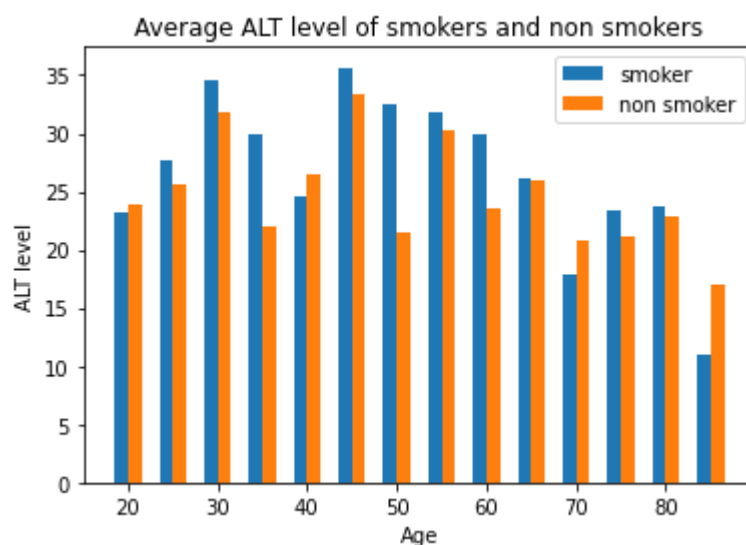
```
In [49]: 1 plt.figure(figsize=(10,10))
2 smokeYes.groupby('age').mean()['LDL'].plot(color='k',marker='o',markerfacecolor='r',
3                                             linestyle='-',linewidth=2)
4 smokeNo.groupby('age').mean()['LDL'].plot(color='k',marker='o',markerfacecolor='g',
5                                             linestyle='-',linewidth=2)
6 plt.xlabel('Age')
7 plt.ylabel("LDL level")
8 plt.title("LDL level of smokers and Non-smokers")
9 plt.legend(loc='lower left')
10 plt.grid()
11 plt.show()
```



Average ALT level of smokers and non smokers in different age.

```
In [50]: 1 plt.figure(figsize=(10,8))
2
3 age=df['age'].unique()
4 altsmokers=smokeYes.groupby('age').mean()['ALT']
5 altnonsmokers=smokeNo.groupby('age').mean()['ALT']
6
7 fig,ax=plt.subplots()
8 bar_width=1.5
9
10 bar_smok_position=[x-bar_width/2 for x in age]
11 bar_nonsmok_position=[x+bar_width/2 for x in age]
12
13 ax.bar(bar_smok_position,altsmokers,bar_width,label='smoker')
14 ax.bar(bar_nonsmok_position,altnonsmokers,bar_width,label='non smoker')
15
16 ax.set_title("Average ALT level of smokers and non smokers")
17 ax.set_xlabel("Age")
18 ax.set_ylabel("ALT level")
19 plt.legend()
20
21 plt.show()
22
23
24
```

<Figure size 720x576 with 0 Axes>




Average ALT level of smokers and non smokers of different age

```

In [51]: 1 plt.figure(figsize=(10,10))
2
3 age=df['age'].unique()
4 astsmokers=smokeYes.groupby('age').mean()['AST']
5 astnonsmokers=smokeNo.groupby('age').mean()['AST']
6
7 fig,ax=plt.subplots()
8 bar_width=1.5
9
10 bar_smok_position=[x-bar_width/2 for x in age]
11 bar_nonsmok_position=[x+bar_width/2 for x in age]
12
13 ax.bar(bar_smok_position,astsmokers,bar_width,label='smoker')
14 ax.bar(bar_nonsmok_position,astnonsmokers,bar_width,label='non smoker')
15
16 ax.set_title("Average AST level of smokers and non smokers")
17 ax.set_xlabel("Age")
18 ax.set_ylabel("AST level")
19 plt.legend()
20
21 plt.show()
22

```

<Figure size 720x720 with 0 Axes>



In []:

1

Statistical test

```

In [52]: 1 df.columns

```

```

Out[52]: Index(['ID', 'smoking', 'gender', 'age', 'height(cm)', 'weight(kg)',
               'waist(cm)', 'eyesight(left)', 'eyesight(right)', 'hearing(left)',
               'hearing(right)', 'systolic', 'systolic_category', 'relaxation',
               'relaxation_category', 'fasting blood sugar', 'diabetic_category',
               'Cholesterol', 'triglyceride', 'HDL', 'LDL', 'hemoglobin',
               'serum creatinine', 'AST', 'ALT', 'Gtp', 'oral', 'dental caries',
               'tartar'],
              dtype='object')

```

To study the association between the smoking status and dental caries

H0: There is no association between the smoking status and fasting sugar level

H1: There is association between the smoking status and fasting sugar level

```
In [53]: 1 crosstab=pd.crosstab(df['smoking'],df['dental caries'])
          2 crosstab
```

```
Out[53]: dental caries    No    Yes
          smoking
          -----
          No    6375   28862
          Yes    5506   14949
```

```
In [54]: 1 import scipy.stats as stats
          2 stats.chi2_contingency(crosstab)
```

```
Out[54]: (600.1881467493236,
          1.5236163632659182e-132,
          1,
          array([[ 7517.25197515, 27719.74802485],
                  [ 4363.74802485, 16091.25197515]]))
```

We reject null hypothesis , because p-value<0.05. Therefore there is a association between the smoking status and fasting blood sugar level.

```
In [55]: 1 df.columns
```

```
Out[55]: Index(['ID', 'smoking', 'gender', 'age', 'height(cm)', 'weight(kg)',
               'waist(cm)', 'eyesight(left)', 'eyesight(right)', 'hearing(left)',
               'hearing(right)', 'systolic', 'systolic_category', 'relaxation',
               'relaxation_category', 'fasting blood sugar', 'diabetic_category',
               'Cholesterol', 'triglyceride', 'HDL', 'LDL', 'hemoglobin',
               'serum creatinine', 'AST', 'ALT', 'Gtp', 'oral', 'dental caries',
               'tartar'],
              dtype='object')
```

Checking whether smoking influences on systolic blood pressure.

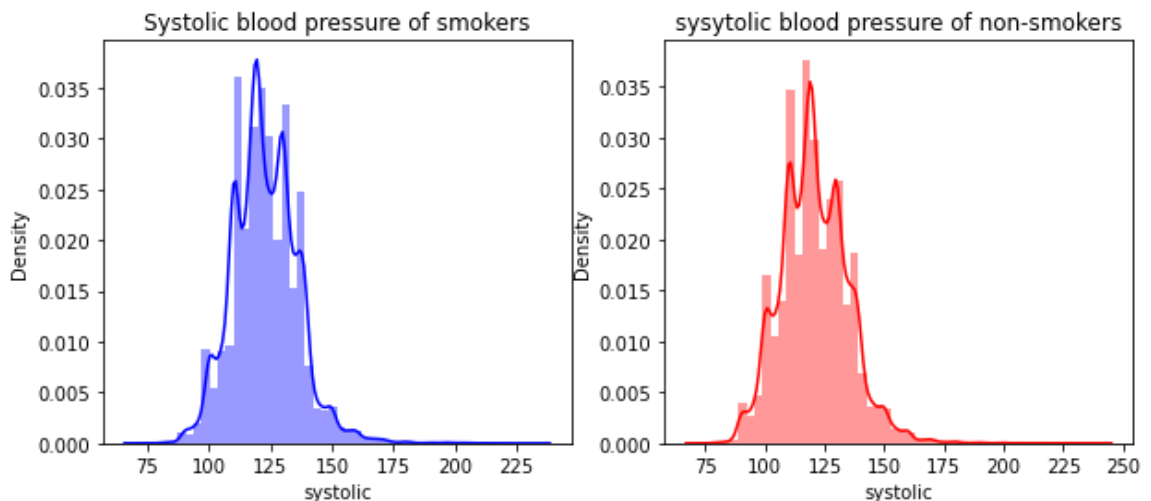
```
In [56]: 1 #Checking for normality of systolic blood pressure
2
3 plt.figure(figsize=(10,4))
4 bp1=smokeyes["systolic"]
5 bp2=smokeno["systolic"]
6
7
8 plt.subplot(1,2,1)
9 sns.distplot(bp1,color='b')
10 plt.title("Systolic blood pressure of smokers")
11
12 plt.subplot(1,2,2)
13 sns.distplot(bp2,color='r')
14 plt.title("systolic blood pressure of non-smokers")
15
16 plt.show()
```

C:\Users\darsh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

C:\Users\darsh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)



H0 : Smoking not influencing on systolic blood pressure of smoker

H1 : Smoking influencing on systolic blood pressure


```
In [57]: 1 from scipy.stats import ttest_ind # t test for relaxation bp
2 t3=smokeyes['systolic']
3 t4=smokeno['systolic']
4 stat,p_value=ttest_ind(t3,t4)
5 print("statistic= ",stat)
6 print("p value= ",p_value)
7
8 alpha = 0.05
9 if p_value < alpha:
10     print("Smoking influesing on systolic blood pressure")
11 else:
12     print("Smoking not influesing on systolic blood pressure of smoker.")
13
14
```

```
statistic= 17.299182876018282
p value= 7.140975601346086e-67
Smoking influesing on systolic blood pressure
```

```
In [ ]: 1
```

```
In [ ]: 1
```

Checking whether smoking influences on relaxation blood pressure.

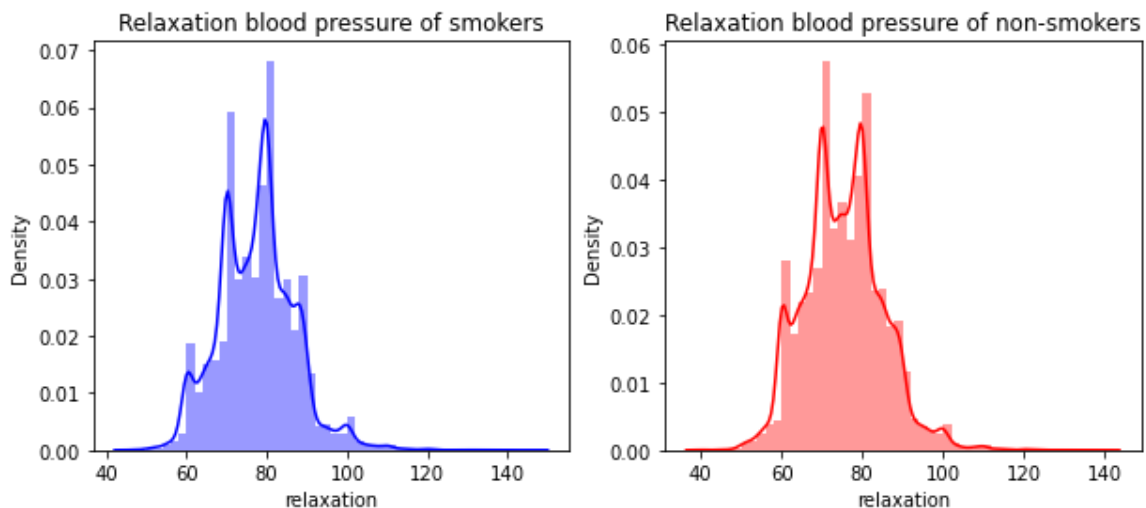
```
In [58]: 1 #Checking for normality of relaxation blood pressure
2 plt.figure(figsize=(10,4))
3 bp3=smokeyes["relaxation"]
4 bp4=smokeno["relaxation"]
5
6
7 plt.subplot(1,2,1)
8 sns.distplot(bp3,color='b')
9 plt.title("Relaxation blood pressure of smokers")
10
11 plt.subplot(1,2,2)
12 sns.distplot(bp4,color='r')
13 plt.title("Relaxation blood pressure of non-smokers")
14
15 plt.show()
```

C:\Users\darsh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

C:\Users\darsh\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)



H0 : Smoking not influencing on relaxation blood pressure of smoker

H1 : Smoking influencing on relaxation blood pressure

```
In [59]: 1 from scipy.stats import ttest_ind # t test for relaxation bp
2 t3=smokeyes['relaxation']
3 t4=smokeno['relaxation']
4 stat,p=ttest_ind(t3,t4)
5 print("statistic= ",stat)
6 print("p value= ",p)
7
8 alpha = 0.05
9 if p_value < alpha:
10     print("Smoking influesing on systolic blood pressure")
11 else:
12     print("Smoking not influesing on systolic blood pressure of smoker.")
13
```

```
statistic= 25.71088809521144
p value= 6.221084836058103e-145
Smoking influesing on systolic blood pressure
```

```
In [ ]: 1
```

```
In [ ]: 1
```

Study the association between the smoking status and fasting blood sugar level

H0: There is no association between the smoking status and fasting sugar level

H1: There is association between the smoking status and fasting sugar level

```
In [60]: 1 crosstab=pd.crosstab(df['smoking'],df['diabetic_category'])
2 crosstab
```

```
Out[60]: diabetic_category  Diabetic  Hypoglycemia  Normal  Pre-Diabetic
smoking
No      1344      2165      22638      9090
Yes     1347      1049      11459      6600
```

```
In [61]: 1 import scipy.stats as stats
2 stats.chi2_contingency(crosstab)
```

```
Out[61]: (564.0462666344756,
6.270104551318755e-122,
3,
array([[ 1702.62815126,  2033.53655821,  21573.58308195,  9927.25220858],
[  988.37184874,  1180.46344179,  12523.41691805,  5762.74779142]]))
```

```
In [ ]: 1
```

Checking whether smokers cholestrol level influencing systolic blood pressure and relaxation blood pressure

In [62]: 1 *#using Levene's method to test the whether variance of group are equal.*

```
In [63]: 1 #Levenes test for systolic blood pressure
2 from scipy.stats import levene
3
4 bp1=smokeyes['Cholesterol'][df['systolic_category']=='Hypotension']
5 bp2=smokeyes['Cholesterol'][df['systolic_category']=='Normal']
6 bp3=smokeyes['Cholesterol'][df['systolic_category']=='Elevated']
7 bp4=smokeyes['Cholesterol'][df['systolic_category']=='Hypertension stage 1']
8 bp5=smokeyes['Cholesterol'][df['systolic_category']=='Hypertension stage 2']
9 bp6=smokeyes['Cholesterol'][df['systolic_category']=='Hypertension crisis']
10
11
12 statistic, p_value = levene(bp1,bp2,bp3,bp4,bp5,bp6)
13
14
15 alpha = 0.05
16
17 print('Statistic :',statistic)
18 print('P-value : ',p_value)
19
20 if p_value < alpha:
21     print("The variances are significantly different (reject the null hypothesis)")
22 else:
23     print("The variances are not significantly different (fail to reject the null hypothesis)")
```

Statistic : 8.035912185426364

P-value : 1.3973650420756164e-07

The variances are significantly different (reject the null hypothesis of equal variances).

```
In [64]: 1 # Levene's test for relaxation blood pressure
2
3 bp1=smokeyes['Cholesterol'][df['relaxation_category']=='Hypotension']
4 bp2=smokeyes['Cholesterol'][df['relaxation_category']=='Normal']
5 bp3=smokeyes['Cholesterol'][df['relaxation_category']=='Elevated']
6 bp4=smokeyes['Cholesterol'][df['relaxation_category']=='Hypertension stage 1']
7 bp5=smokeyes['Cholesterol'][df['relaxation_category']=='Hypertension stage 2']
8 bp6=smokeyes['Cholesterol'][df['relaxation_category']=='Hypertension crisis']
9
10
11 statistic, p_value = levene(bp1,bp2,bp3,bp4,bp5,bp6)
12
13
14 alpha = 0.05
15
16 print('Statistic :',statistic)
17 print('P-value : ',p_value)
18
19 if p_value < alpha:
20     print("The variances are significantly different (reject the null hypothesis)")
21 else:
22     print("The variances are not significantly different (fail to reject the null hypothesis)")
```

Statistic : 4.950484091083053

P-value : 0.00015644784666161838

The variances are significantly different (reject the null hypothesis of equal variances).

Using kruskal-wallis test checking whether smokers cholestrol level influencing systolic blood pressure

Ho: cholestrol level do not influencing systolic blood pressure

H1: cholestrol level influencing systolic blood pressure

```
In [65]: 1 import scipy.stats as stats
2
3 bp1=smokeyes['Cholesterol'][df['systolic_category']=='Hypotension']
4 bp2=smokeyes['Cholesterol'][df['systolic_category']=='Normal']
5 bp3=smokeyes['Cholesterol'][df['systolic_category']=='Elevated']
6 bp4=smokeyes['Cholesterol'][df['systolic_category']=='Hypertension stage 1']
7 bp5=smokeyes['Cholesterol'][df['systolic_category']=='Hypertension stage 2']
8 bp6=smokeyes['Cholesterol'][df['systolic_category']=='Hypertension crisis']
9
10 # Perform the Kruskal-Wallis test
11 statistic, p_value = stats.kruskal(bp1,bp2,bp3,bp4,bp5,bp6)
12
13
14 print("Kruskal-Wallis Test")
15 print("Statistic:", statistic)
16 print("P-value:", p_value)
17
18 # Check the significance level and draw conclusions
19 alpha = 0.05
20 if p_value < alpha:
21     print("There is a significant difference between the groups.")
22 else:
23     print("There is no significant difference between the groups.")
```

Kruskal-Wallis Test

Statistic: 112.85332939360227

P-value: 1.0216312842082404e-22

There is a significant difference between the groups.

In []:

1

In []:

1

Checking whether smokers cholestrol level influencing relaxation blood pressure

Ho: cholestrol level do not influencing systolic blood pressure

H1: cholestrol level influencing systolic blood pressure

```
In [66]: 1
2 bp1=smokeyes['Cholesterol'][df['relaxation_category']=='Hypotension']
3 bp2=smokeyes['Cholesterol'][df['relaxation_category']=='Normal']
4 bp3=smokeyes['Cholesterol'][df['relaxation_category']=='Elevated']
5 bp4=smokeyes['Cholesterol'][df['relaxation_category']=='Hypertension sta
6 bp5=smokeyes['Cholesterol'][df['relaxation_category']=='Hypertension sta
7 bp6=smokeyes['Cholesterol'][df['relaxation_category']=='Hypertension cr
8
9 # Perform the Kruskal-Wallis test
10 statistic, p_value = stats.kruskal(bp1,bp2,bp3,bp4,bp5,bp6)
11
12
13 print("Kruskal-Wallis Test")
14 print("Statistic:", statistic)
15 print("P-value:", p_value)
16
17 # Check the significance level and draw conclusions
18 alpha = 0.05
19 if p_value < alpha:
20     print("There is a significant difference between the groups.")
21 else:
22     print("There is no significant difference between the groups.")
```

Kruskal-Wallis Test

Statistic: 210.87431690442494

P-value: 1.3373403109815881e-43

There is a significant difference between the groups.

```
In [67]: 1 df['diabetic_category'].unique()
```

```
Out[67]: array(['Normal', 'Pre-Diabetic', 'Hypoglycemia', 'Diabetic'], dtype=object)
```

Checking whether smokers diabetic level influencing on serum creatinine

```
In [68]: 1 # Levene's test for relaxation blood pressure
2
3 from scipy.stats import levene
4
5 sc1=smokeyes['serum creatinine'][df['diabetic_category']=='Hypoglycemia']
6 sc2=smokeyes['serum creatinine'][df['diabetic_category']=='Normal']
7 sc3=smokeyes['serum creatinine'][df['diabetic_category']=='Pre-Diabetic']
8 sc4=smokeyes['serum creatinine'][df['diabetic_category']=='Diabetic']
9
10
11
12 statistic, p_value = levene(sc1,sc2,sc3,sc4)
13
14
15 alpha = 0.05
16
17 print('Statistic :',statistic)
18 print('P-value : ',p_value)
19
20 if p_value < alpha:
21     print("The variances are significantly different (reject the null hypothesis)")
22 else:
23     print("The variances are not significantly different (fail to reject the null hypothesis)")
```

Statistic : 9.784655477556692

P-value : 1.9055390702653126e-06

The variances are significantly different (reject the null hypothesis of equal variances).

Kruskal-Wallis Test for checking whether smokers diabetic level influencing on serum creatine.

```

In [69]: 1 from scipy.stats import kruskal
2
3 sc1=smokeyes['serum creatinine'][df['diabetic_category']=='Hypoglycemia']
4 sc2=smokeyes['serum creatinine'][df['diabetic_category']=='Normal']
5 sc3=smokeyes['serum creatinine'][df['diabetic_category']=='Pre-Diabetic']
6 sc4=smokeyes['serum creatinine'][df['diabetic_category']=='Diabetic']
7
8
9
10 statistic, p_value = kruskal(sc1,sc2,sc3,sc4)
11
12 print("Kruskal-Wallis Test")
13 print("Statistic:", statistic)
14 print("P-value:", p_value)
15
16 # Check the significance level and draw conclusions
17 alpha = 0.05
18 if p_value < alpha:
19     print("There is a significant difference between serum creatinine of")
20 else:
21     print("There is no significant difference between serum creatinine of")
22
23

```

Kruskal-Wallis Test

Statistic: 17.37610751015689

P-value: 0.0005913810964908272

There is a significant difference between serum creatinine of smokers .

In []:

1

In []:

1

In [70]:

1 #Model building

In [71]:

1 df.head(3)

Out[71]:


	ID	smoking	gender	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
0	0	No	F	40	155	60	81.3	1.2	1.0
1	1	No	F	40	160	60	81.0	0.8	0.6
2	2	Yes	M	55	170	60	80.0	0.8	0.8

3 rows × 29 columns




```
In [72]: 1 plt.figure(figsize=(40,40))
2         sns.heatmap(df.corr(),annot=True,cmap='RdYlGn',annot_kws={'size':15})
```

Out[72]: <AxesSubplot:>



```
In [73]: 1 from sklearn.metrics import classification_report,PrecisionRecallDisplay
2         from sklearn.metrics import precision_score,recall_score,f1_score
```

```
In [74]: 1 from sklearn.preprocessing import LabelEncoder
2         for column in df.columns:
3             if df[column].dtype==np.number:
4                 continue
5             else:
6                 df[column]=LabelEncoder().fit_transform(df[column])
```

C:\Users\darsh\AppData\Local\Temp\ipykernel_18136\4259945406.py:3: Deprecat
ionWarning: Converting `np.inexact` or `np.floating` to a dtype is deprecate
d. The current result is `float64` which is not strictly correct.
if df[column].dtype==np.number:

In [75]: 1 df.columns

Out[75]: Index(['ID', 'smoking', 'gender', 'age', 'height(cm)', 'weight(kg)', 'waist(cm)', 'eyesight(left)', 'eyesight(right)', 'hearing(left)', 'hearing(right)', 'systolic', 'systolic_category', 'relaxation', 'relaxation_category', 'fasting blood sugar', 'diabetic_category', 'Cholesterol', 'triglyceride', 'HDL', 'LDL', 'hemoglobin', 'serum creatinine', 'AST', 'ALT', 'Gtp', 'oral', 'dental caries', 'tartar'], dtype='object')

In [76]: 1 x=df.drop(columns=['ID','eyesight(left)','eyesight(right)','hearing(left)'])
2 y=df['diabetic_category']

In [77]: 1 from sklearn.model_selection import train_test_split
2 xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.20,random_state=42)

LogisticRegression

In [78]: 1 from sklearn.linear_model import LogisticRegression
2 model=LogisticRegression()
3 model.fit(xtrain,ytrain)

C:\Users\darsh\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n_iter_i = _check_optimize_result(

Out[78]: LogisticRegression()

In [79]: 1 ypred1=model.predict(xtest)

In [80]: 1 from sklearn.metrics import accuracy_score
2 print("Test Accuracy: ",accuracy_score(ytest,ypred1))
3 print("Precision: ",precision_score(ytest,ypred1,average='weighted'))

Test Accuracy: 0.625100996498788
Precision: 0.5706326704136191

In [81]: 1 ypredd=model.predict(xtrain)

In [82]: 1 from sklearn.metrics import accuracy_score
2 print("Train Accuracy: ",accuracy_score(ytrain,ypredd))

Train Accuracy: 0.6209682849639755

```
In [83]: 1 from sklearn.metrics import classification_report
2 print(classification_report(ytest,ypred1))
```

	precision	recall	f1-score	support
0	0.27	0.01	0.02	493
1	0.50	0.00	0.00	653
2	0.65	0.92	0.76	6869
3	0.46	0.20	0.28	3124
accuracy			0.63	11139
macro avg	0.47	0.28	0.26	11139
weighted avg	0.57	0.63	0.55	11139

```
In [84]: 1 print()
```

KNN

```
In [85]: 1 from sklearn.neighbors import KNeighborsClassifier
2 from sklearn.metrics import f1_score
```

```
In [86]: 1 clf=KNeighborsClassifier(n_neighbors=5)
2 clf.fit(xtrain,ytrain)
3 ypred2=clf.predict(xtest)
```

```
In [87]: 1 from sklearn.metrics import accuracy_score
2 print("Test Accuracy: ",accuracy_score(ytest,ypred2))
```

Test Accuracy: 0.5819193823503007

```
In [88]: 1 ypred2=clf.predict(xtrain)
```

```
In [89]: 1 from sklearn.metrics import accuracy_score
2 print("Train Accuracy: ",accuracy_score(ytrain,ypred2))
```

Train Accuracy: 0.7083249163917132

```
In [90]: 1 print(classification_report(ytest,ypred2))
```

	precision	recall	f1-score	support
0	0.18	0.14	0.15	493
1	0.14	0.08	0.10	653
2	0.66	0.79	0.72	6869
3	0.43	0.29	0.35	3124
accuracy			0.58	11139
macro avg	0.35	0.32	0.33	11139
weighted avg	0.54	0.58	0.55	11139

In []:

1

In [91]:

```

1 from sklearn.neighbors import KNeighborsClassifier
2 def elbow(k):
3     test_error=[]
4     for i in k:
5         clf=KNeighborsClassifier(n_neighbors=i)
6         clf.fit(xtrain,ytrain)
7         tmp=clf.predict(xtest)
8         tmp=f1_score(ytest,tmp,average='micro')
9         error=1-tmp
10        test_error.append(error)
11    return test_error

```

In [92]:

```

1 k=range(6,20)
2 test=elbow(k)
3 print(test)

```

```

[0.40506329113924056, 0.399138163210342, 0.3935721339438011, 0.398419965885
6271, 0.3907891193105306, 0.38800610467726004, 0.38764700601490254, 0.38971
18233234581, 0.3868390340245983, 0.3847742167160427, 0.383876470060149, 0.3
854026393751683, 0.3837866953945597, 0.3846844420504534]

```


In [93]:

```

1 plt.plot(k,test)
2 plt.xlabel("K Neighbors")
3 plt.ylabel("test_error")
4 plt.title("Elbow curve for test")

```

Out[93]: Text(0.5, 1.0, 'Elbow curve for test')



In [94]:

```

1 knn=KNeighborsClassifier(n_neighbors=16)
2 knn.fit(xtrain,ytrain)

```

Out[94]: KNeighborsClassifier(n_neighbors=16)

In [95]:

```

1 ytrain_prdct=knn.predict(xtrain)
2 ytest_prdct=knn.predict(xtest)

```

```
In [96]: 1 print("Test Accuracy :",accuracy_score(ytest_prdct,ytest))
          2 print("Train Accuracy :",accuracy_score(ytrain_prdct,ytrain))
```

Test Accuracy : 0.616123529939851
Train Accuracy : 0.6430094494197922

```
In [ ]: 1
```

Naive Bayesian Classification

```
In [97]: 1 from sklearn.naive_bayes import GaussianNB
          2 gnb=GaussianNB()
          3 gnb.fit(xtrain,ytrain)
```

Out[97]: GaussianNB()

```
In [98]: 1 ypred3=gnb.predict(xtest)
          2 ypred3=gnb.predict(xtrain)
```

```
In [99]: 1 from sklearn.metrics import accuracy_score
          2 print("Test Accuracy: ",accuracy_score(ytest,ypred3))
          3 print("Train Accuracy: ",accuracy_score(ytrain,ypred3))
```

Test Accuracy: 0.5688122811742526
Train Accuracy: 0.5723969205216259

```
In [100]: 1 print(classification_report(ytest,ypred3))
```

	precision	recall	f1-score	support
0	0.17	0.21	0.19	493
1	0.10	0.03	0.04	653
2	0.67	0.77	0.71	6869
3	0.39	0.31	0.34	3124
accuracy			0.57	11139
macro avg	0.33	0.33	0.32	11139
weighted avg	0.53	0.57	0.55	11139

RandomForest and DecisionTreeClassifier

```
In [101]: 1 from sklearn.tree import DecisionTreeClassifier
          2 dtree=DecisionTreeClassifier()
          3 dtree.fit(xtrain,ytrain)
```

Out[101]: DecisionTreeClassifier()

```
In [102]: 1 ypred4=dtree.predict(xtest)
          2 ypred4=dtree.predict(xtrain)
```

```
In [103]: 1 print("Test Accuracy: ",accuracy_score(ytest,ypred4))
          2 print("Train Accuracy: ",accuracy_score(ytrain,ypredd4))
```

Test Accuracy: 0.6531106921626717
Train Accuracy: 1.0

```
In [104]: 1 print(classification_report(ytest,ypred4))
```

	precision	recall	f1-score	support
0	0.35	0.42	0.38	493
1	0.35	0.37	0.36	653
2	0.76	0.75	0.75	6869
3	0.54	0.55	0.54	3124
accuracy			0.65	11139
macro avg	0.50	0.52	0.51	11139
weighted avg	0.66	0.65	0.66	11139

Random Forest

```
In [105]: 1 from sklearn.ensemble import RandomForestClassifier
          2 rf=RandomForestClassifier()
          3 rf.fit(xtrain,ytrain)
```

Out[105]: RandomForestClassifier()

```
In [106]: 1 ypred5=rf.predict(xtest)
          2 ypred5=rf.predict(xtrain)
```

```
In [107]: 1 print("Test Accuracy: ",accuracy_score(ytest,ypred5))
          2 print("Train Accuracy: ",accuracy_score(ytrain,ypredd5))
          3
```

Test Accuracy: 0.7467456683723853
Train Accuracy: 1.0

```
In [108]: 1 # hyperparameter tuning
          2
```

GridSearchCV

```
In [110]: 1 from sklearn.model_selection import GridSearchCV
          2 from sklearn.model_selection import RandomizedSearchCV
```

```
In [111]: 1 param_grid = {
          2     'n_estimators': [25, 50, 100, 150],
          3     'max_features': ['sqrt', 'log2', None],
          4     'max_depth': [3, 6, 9],
          5     'max_leaf_nodes': [3, 6, 9],
          6 }
```

```
In [113]: 1 grid_search = GridSearchCV(RandomForestClassifier(),
2                                param_grid=param_grid)
3 grid_search.fit(xtrain, ytrain)
4 print(grid_search.best_estimator_)
```

RandomForestClassifier(max_depth=9, max_features=None, max_leaf_nodes=9, n_estimators=150)

```
In [114]: 1 model_grid = RandomForestClassifier(max_depth=9,
2                                           max_features=None,
3                                           max_leaf_nodes=9,
4                                           n_estimators=150)
5 model_grid.fit(xtrain, ytrain)
6 y_pred_grid = model.predict(xtest)
7 print(classification_report(y_pred_grid, ytest))
```

	precision	recall	f1-score	support
0	0.01	0.27	0.02	15
1	0.00	0.50	0.00	2
2	0.92	0.65	0.76	9770
3	0.20	0.46	0.28	1352
accuracy			0.63	11139
macro avg	0.28	0.47	0.26	11139
weighted avg	0.83	0.63	0.70	11139

```
In [116]: 1 print("Test Accuracy: ",accuracy_score(y_pred_grid, ytest))
```

Test Accuracy: 0.625100996498788

```
In [118]: 1 ytrain_pred_grid = model.predict(xtrain)
2 print("Train Accuracy",accuracy_score(ytrain_pred_grid, ytrain))
```

Train Accuracy 0.6209682849639755

```
In [119]: 1 # Randomizedsearch
```

```
In [125]: 1 param_grid = {
2     'n_estimators': [25, 50, 100, 150],
3     'max_features': ['sqrt', 'log2', None],
4     'max_depth': [9, 12, 15],
5     'max_leaf_nodes': [9, 12, 15],
6 }
```

```
In [126]: 1 random_search = RandomizedSearchCV(RandomForestClassifier(),
2                                       param_grid)
3 random_search.fit(xtrain, ytrain)
4 print(random_search.best_estimator_)
```

RandomForestClassifier(max_depth=15, max_features=None, max_leaf_nodes=15, n_estimators=25)

```
In [128]: 1 model_random = RandomForestClassifier(max_depth=15,
2                                           max_features=None,
3                                           max_leaf_nodes=15,
4                                           n_estimators=25)
5 model_random.fit(xtrain, ytrain)
6 ytest_pred_rand1 = model.predict(xtest)
7 ytrain_pred_rand1=model.predict(xtrain)
8 print("Test Accuracy :",accuracy_score(ytest_pred_rand1, ytest))
9 print("Train Accuracy :",accuracy_score(ytrain_pred_rand1,ytrain))
```

Test Accuracy : 0.625100996498788

Train Accuracy : 0.6209682849639755

Xgboost

```
In [ ]: 1 params={
2       'objective':'multi:softmax',
3       'num_class':len(set(y)),
4       'max_depth':3,
5       'learning_rate':0.1,
6       'n_estimators':100
7       }
```

```
In [134]: 1 !pip install xgboost
```

Collecting xgboost

Downloading xgboost-1.7.6-py3-none-win_amd64.whl (70.9 MB)

Requirement already satisfied: numpy in c:\users\darsh\anaconda3\lib\site-packages (from xgboost) (1.20.3)

Requirement already satisfied: scipy in c:\users\darsh\anaconda3\lib\site-packages (from xgboost) (1.7.1)

Installing collected packages: xgboost

Successfully installed xgboost-1.7.6

```
In [135]: 1 import xgboost as xgb
2
```

```
In [137]: 1 model=xgb.XGBClassifier()
2
3 model.fit(xtrain,ytrain)
4 ytest_pred=model.predict(xtest)
5 ytrain_pred=model.predict(xtrain)
6
7 print("Test ACcuracy: ",accuracy_score(ytest_pred,ytest))
8 print("Train Accuracy: ",accuracy_score(ytrain_pred,ytrain))
```

Test ACcuracy: 0.6548164108088698

Train Accuracy: 0.750723856979328

```
In [ ]: 1
```



```
In [138]: 1 booster=['gbtree', 'gblinear']
2 base_score=[0.25,0.5,0.75,1]
3 n_estimators = [100, 500, 900, 1100, 1500]
4 max_depth = [2, 3, 5, 10, 15]
5 booster=['gbtree', 'gblinear']
6 learning_rate=[0.05,0.1,0.15,0.20]
7 min_child_weight=[1,2,3,4]
8
9 # Define the grid of hyperparameters to search
10 hyperparameter_grid = {
11     'n_estimators': n_estimators,
12     'max_depth':max_depth,
13     'learning_rate':learning_rate,
14     'min_child_weight':min_child_weight,
15     'booster':booster,
16     'base_score':base_score
17 }
```

```
In [140]: 1 random_cv = RandomizedSearchCV(estimator=xgb.XGBClassifier(),
2     param_distributions=hyperparameter_grid,
3     cv=5, n_iter=50,
4     scoring = 'neg_mean_absolute_error',n_jobs = 4,
5     verbose = 5,
6     return_train_score = True,
7     random_state=42)
```

```
In [ ]: 1 random_cv.fit(xtrain,ytrain)
2
3 random_cv.best_estimator_
4
```

```
In [ ]: 1
```